



Fault-tolerant TopK algorithms

Lewis Tseng
Boston College

Ruifan Yang
Cornell University

Zheng Zhou
Boston College

- In elections:

Which are the k globally most popular candidates?

- In peer-to-peer file-sharing networks:
which are most popular downloads?

Problem Statement

A top-k Query :

Across all nodes in the system,

return k objects with highest global attribute (in this case $N_i = \sum_j A_{ij}$).

2-D array

n items, m peers

Each item i has a local attribute

in each peer j, A_{ij}



Threshold Algorithm : Book a hotel

Step 1:

Parallel sorted access to each list/peer

For each object seen:

- Get all grades by random access
- Determine $\text{Min}(A_1, A_2)$
- Keep the current highest 2 in buffer

Rate	Distance
(a, 0.9)	(d, 0.9)
(b, 0.8)	(a, 0.85)
(c, 0.72)	(b, 0.7)
⋮	⋮
(d, 0.6)	(c, 0.2)

ID	A_1	A_2	$\text{Min}(A_1, A_2)$
a	0.9	0.85	0.85
d	0.6	0.9	0.6

Threshold Algorithm Cont.

Step 2:

Determine threshold value T based on objects currently seen under sorted access

$$T = \min(L1, L2)$$

If 2 objects with overall grade \geq threshold value
Stop

else

go to next entry position in sorted list
& repeat step 1

Rate

Distance



(a, 0.9)

(d, 0.9)

(b, 0.8)

(a, 0.85)

(c, 0.72)

(b, 0.7)

.

.

.

.

.

.

(d, 0.6)

(c, 0.2)

$$T = \min(0.9, 0.9) = 0.9$$

ID	A ₁	A ₂	Min(A ₁ , A ₂)
a	0.9	0.85	0.85
d	0.6	0.9	0.6

Threshold Algorithm Cont.

Step 1 Again

Parallel sorted access to each list/peer

For each object seen:

- Get all grades by random access
- Determine $\text{Min}(A_1, A_2)$
- Keep the current highest 2 in buffer

Rate	Distance
(a, 0.9)	(d, 0.9)
(b, 0.8)	(a, 0.85)
(c, 0.72)	(b, 0.7)
.	.
.	.
.	.
(d, 0.6)	(c, 0.2)

ID	A_1	A_2	$\text{Min}(A_1, A_2)$
a	0.9	0.85	0.85
d	0.6	0.9	0.6
b	0.8	0.7	0.7

Threshold Algorithm Cont.

Step 2 Again

Determine threshold value T based on objects currently seen under sorted access

$$T = \min(L1, L2)$$

If 2 objects with overall grade \geq threshold value
Stop

else
go to next entry position in sorted list
& repeat step 1



Rate

(a, 0.9)

(b, 0.8)

(c, 0.72)

.

.

.

(d, 0.6)

Distance

(d, 0.9)

(a, 0.85)

(b, 0.7)

.

.

.

(c, 0.2)

$$T = \min(0.8, 0.85) = 0.8$$

ID	A_1	A_2	$\text{Min}(A_1, A_2)$
a	0.9	0.85	0.85
b	0.8	0.7	0.7

Threshold Algorithm Cont.

Stop

Here

Rate

(a, 0.9)

(b, 0.8)

(c, 0.72)

.

.

.

(d, 0.6)

Distance

(d, 0.9)

(a, 0.85)

(b, 0.7)

.

.

.

(c, 0.2)

$$T = \min(0.72, 0.7) = 0.7$$

ID	A ₁	A ₂	Min(A ₁ , A ₂)
a	0.9	0.85	0.85
b	0.8	0.7	0.7

Definition: Safe Algorithm

An algorithm is **safe** if the algorithm will never output an item with no attribute in non-faulty nodes.

What if there are faulty nodes ?

Ex: Three peers &
one faulty node

	A	B	C
item 1	0	0	1
item 2	0	1	0
item 3	1	0	0

**All safe algorithms
perform badly here.**

Three Faulty Behaviors

(1) Zero: The faulty nodes set all the attributes in it to zero.

(2) Uniform: The faulty nodes set all the attributes in it to a random number between 0 and 10 times the maximum value of all attributes.

(3) Large: The faulty nodes set all the attributes in it to a random number between 5 and 10 times the maximum value of all attributes.

Unsafe Algorithms

Algorithm 1: In the beginning of the algorithm, we randomly picked f peers and assume them to be faulty. Then run the Threshold Algorithm.

Algorithm 2: Randomly pick one node to exclude from our system for this round. Then, run the Threshold Algorithm to find the top one item among the remaining peer, and add it to the output. Then randomly pick another node. Run the Threshold Algorithm for the remaining peers. Add the highest ranking item that's not in the output yet. Repeat until we have k items.


Algorithm 3: We first find the median of each item over all peers, and then give rank of each item according to their median.

Experiments


1. Consider all three faulty behaviors
2. $k \in [1,3,5,7,10,15]$ (that is, different length of output)
3. For each k , make 500 datasets with $\text{zipf}(1.34/1.3/1.25, 50)$
4. Run the algorithms in the 500 datasets
5. Calculate the average accuracy

Results :

Approaching to 1: accuracy increases






Zipf(1.34, 50) with zero faulty nodes



Zipf(1.25, 50) with zero faulty nodes

Results :

For zipfian: zero nodes decrease the accuracy the most



Zero nodes

Results

For Uniformly distributed datasets

